

CLD Lab 02: App scaling on IaaS

Authors: Kevin Auberson and Léo Zmoos - L2GrR

Date: April 7, 2024

TASK 1: CREATE A DATABASE USING THE RELATIONAL DATABASE SERVICE (RDS)

Copy the estimated monthly cost for the database and add it to your report.

RDS Pricing:

- db.t3.micro = \$0.017/h
- gp2 storage = \$0.115/GB-month

Monthly cost calculation:

- 20Gb gp2 storage = $20 * 0.115 = 2.3\$$
- db.t3.micro = $24 * 30 * 0.017 = 12.24\$$

Total monthly cost = 2.3 + 12.24 = 14.54\$

Compare the costs of your RDS instance to a continuously running EC2 instance of the same instance type to see how much AWS charges for the extra functionality.

RDS monthly cost = 14.54\$

EC2 Pricing:

| Instance name ▲ | On-Demand hourly rate ▼ | vCPU ▼ | Memory ▼ | Storage ▼ | Network performance ▼ |
|-----------------|-------------------------|--------|----------|-----------|-----------------------|
| t3.micro | \$0.0104 | 2 | 1 GiB | EBS Only | Up to 5 Gigabit |

Instance total cost: $24 * 30 * 0.0104 = 7,488\$$

GP2 Storage: 0.10\$ /GB-Month => total storage cost = $20 * 0.10 = 2\$$

EC2 total monthly cost = $7,488 + 2 = 9,488\$$

Difference between RDS and EC2: $14.54 - 9.488 = 5,052\$$

They charge an extra ~5\$ for the extra fonctionnality.

In a two-tier architecture the web application and the database are kept separate and run on different hosts. Imagine that for the second tier instead of using RDS to store the data you would create a virtual machine in EC2 and install and run yourself a database on it. If you were the Head of IT of a medium-size business, how would you argue in favor of using a database as a service instead of running your own database on an EC2 instance? How would you argue against it?

Arguments in favor of RDS:

RDS provides automatic backups and encryption, taking full responsibility for the database's configuration, management, maintenance, and security through AWS automation. This enables users to configure read replicas or set up synchronous replication across multiple AZs for improved performance and availability. While deploying a database in multiple AZs with multiple standbys can be time-consuming, it can be a huge time saver with RDS. However, time is money, and this solution is more expensive.

Arguments in favor of EC2 with DB engine

It gives more flexibility and granularity on various aspects of the system. For instance, users can choose a specific OS running a specific version, or they could use EBS RAID and stripping configurations for higher performance. EC2 would be the only solution if a user wants to run a DB engine in an older unsupported version, or if data access time and bandwidth are critical. It could also be a cheaper choice for test/dev DB environments that do not need to be in production.

In conclusion, the decision between RDS and EC2 with a DB engine heavily depends on the required flexibility of the database, the nature of the data stored, and the performance needed for the application.

Copy the endpoint address of the database into the report.

`grr-zmoos-wordpress-db.crsk2uw660uh.us-east-1.rds.amazonaws.com` (`http://grr-zmoos-wordpress-db.crsk2uw660uh.us-east-1.rds.amazonaws.com`)

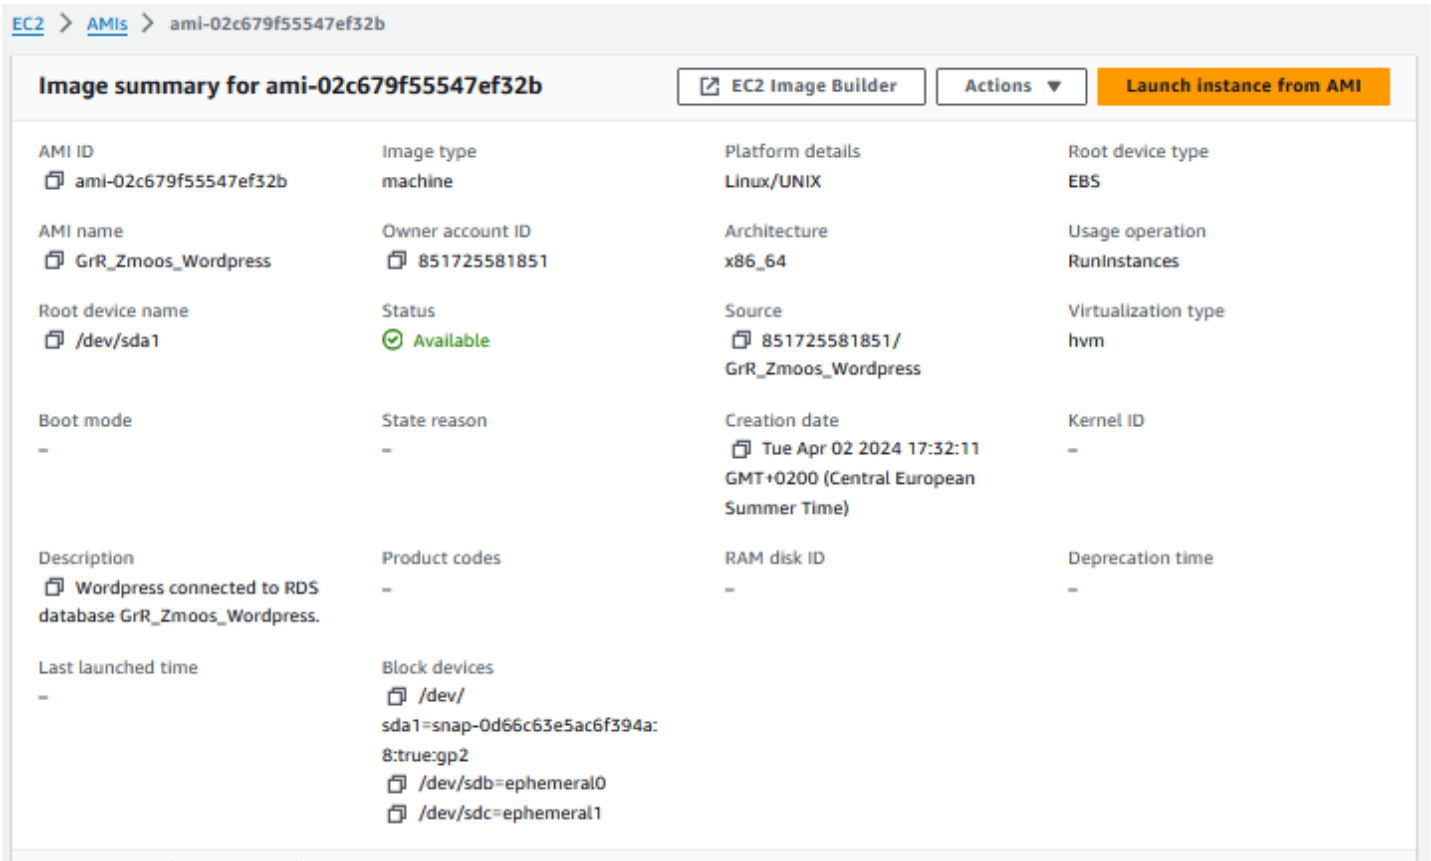
TASK 2: CONFIGURE THE WORDPRESS MASTER INSTANCE TO USE THE RDS DATABASE

Copy the part of /var/www/html/wp-config.php that configures the database into the report.

```
// ** Database settings - You can get this info from your web host ** //  
/** The name of the database for WordPress */  
define( 'DB_NAME', 'grr-zmoos-wordpress-db' );  
  
/** Database username */  
define( 'DB_USER', 'admin' );  
  
/** Database password */  
define( 'DB_PASSWORD', 'dCU=U!dZjqB)2)-' );  
  
/** Database hostname */  
define( 'DB_HOST', 'grr-zmoos-wordpress-db.crsk2uw660uh.us-east-1.rds.amazonaws.com' );
```

TASK 3: CREATE A CUSTOM VIRTUAL MACHINE IMAGE

Copy a screenshot of the AWS console showing the AMI parameters into the report.

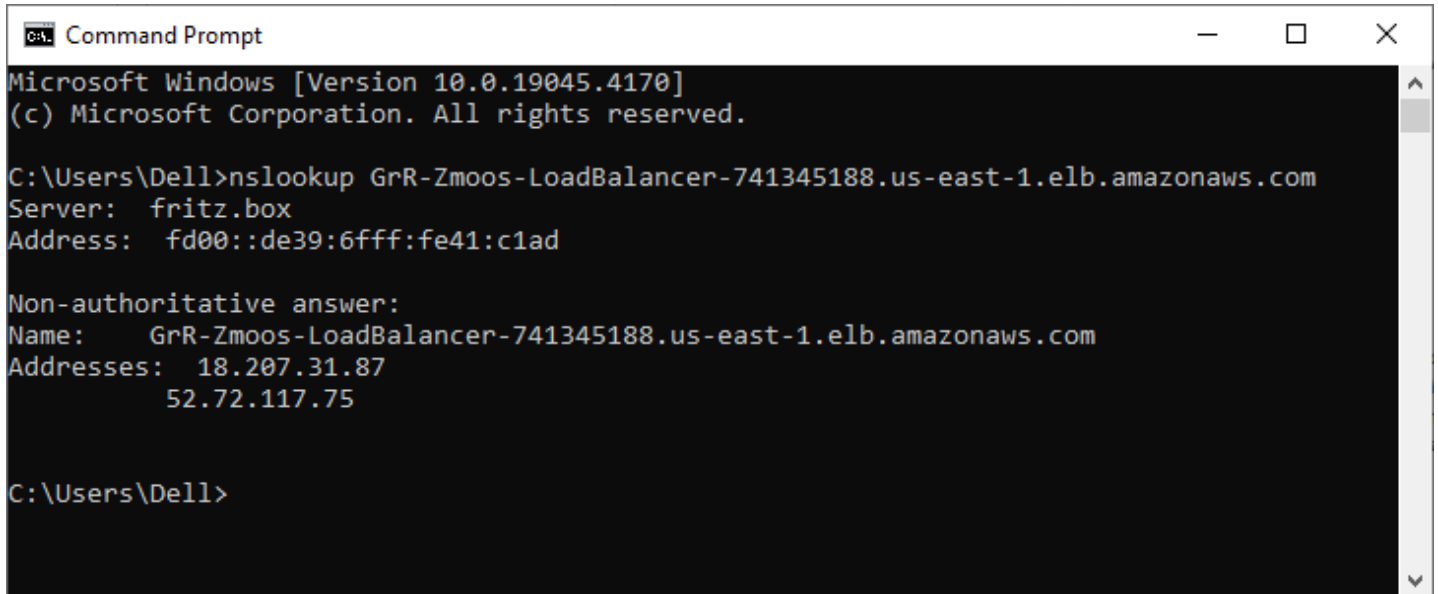


AMI parameters

TASK 4: CREATE A LOAD BALANCER

On your local machine resolve the DNS name of the load balancer into an IP address using the nslookup command (works on Linux, macOS and Windows). Write the DNS name and the resolved IP Address(es) into the report.

DNS Name: GrR-Zmoos-LoadBalancer-741345188.us-east-1.elb.amazonaws.com
(<http://GrR-Zmoos-LoadBalancer-741345188.us-east-1.elb.amazonaws.com>)



```
Command Prompt
Microsoft Windows [Version 10.0.19045.4170]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Dell>nslookup GrR-Zmoos-LoadBalancer-741345188.us-east-1.elb.amazonaws.com
Server: fritz.box
Address: fd00::de39:6fff:fe41:c1ad

Non-authoritative answer:
Name: GrR-Zmoos-LoadBalancer-741345188.us-east-1.elb.amazonaws.com
Addresses: 18.207.31.87
           52.72.117.75

C:\Users\Dell>
```

Screenshot of the nslookup result

In the Apache access log identify the health check accesses from the load balancer and copy some samples into the report.

```

PuTTY (inactive)
172.31.11.220 - - [03/Apr/2024:18:03:06 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.29.116 - - [03/Apr/2024:18:03:15 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.11.220 - - [03/Apr/2024:18:03:16 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.29.116 - - [03/Apr/2024:18:03:25 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.11.220 - - [03/Apr/2024:18:03:26 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.29.116 - - [03/Apr/2024:18:03:35 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.11.220 - - [03/Apr/2024:18:03:36 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.29.116 - - [03/Apr/2024:18:03:45 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.11.220 - - [03/Apr/2024:18:03:46 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.29.116 - - [03/Apr/2024:18:03:55 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.11.220 - - [03/Apr/2024:18:03:56 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.29.116 - - [03/Apr/2024:18:04:05 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"
172.31.11.220 - - [03/Apr/2024:18:04:06 +0000] "GET / HTTP/1.1" 200 3423 "-" "EL
B-HealthChecker/2.0"

```

Log with HealthChecker

TASK 5: LAUNCH A SECOND INSTANCE FROM THE CUSTOM IMAGE

Using the custom virtual machine image you created earlier launch a second instance.

Launch an instance Info

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags Info

Name

GrR_Auberson_Image

Add additional tags

First, we choose a name for the instance.

▼ **Application and OS Images (Amazon Machine Image)** Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

🔍 Search our full catalog including 1000s of application and OS images

Recents

My AMIs

Quick Start

☒ Owned by me

☐ Shared with me

🔍

[Browse more AMIs](#)

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

GrR_Auberson_Wordpress

ami-0c97c41f245bd6816

2024-04-08T14:05:55.000Z Virtualization: hvm ENA enabled: true Root device type: ebs

Description

Wordpress connected to RDS database GrR-Auberson-Wordpress-DB

Architecture

AMI ID

x86_64

ami-0c97c41f245bd6816

Then, the OS Images in this case the image we generate at Task 4.

▼ Instance type [Info](#) | [Get advice](#)

Instance type

t2.micro

Free tier eligible

Family: t2 1 vCPU 1 GiB Memory Current generation: true

On-Demand Windows base pricing: 0.0162 USD per Hour

On-Demand SUSE base pricing: 0.0116 USD per Hour

On-Demand RHEL base pricing: 0.0716 USD per Hour

On-Demand Linux base pricing: 0.0116 USD per Hour

☐ All generations

[Compare instance types](#)

[Additional costs apply for AMIs with pre-installed software](#)

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

GrR_Auberson

 [Create new key pair](#)

The instance type and the key pair for logging, we use the same as the lab 1.

▼ Network settings [Info](#)

[Edit](#)

Network [Info](#)

vpc-049e2f8e56e0bafef

Subnet [Info](#)

No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)

Enable

[Additional charges apply](#) when outside of [free tier allowance](#)

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☐ Create security group


☒ Select existing security group

Common security groups [Info](#)

Select security groups ▼

GrR_Auberson_default sg-02540dd46a8269193 ✕

VPC: vpc-049e2f8e56e0bafef

 [Compare security group rules](#)

Security groups that you add or remove here will be added to or removed from all your network interfaces.

► Configure storage [Info](#)

[Advanced](#)

For network, we use the same security group as the other instance.

▼ Summary

Number of instances

[Info](#)

1

Software Image (AMI)

Wordpress connected to RDS dat...[read more](#)

ami-0c97c41f245bd6816

Virtual server type (instance type)

t2.micro

Firewall (security group)

GrR_Auberson_default

Storage (volumes)

1 volume(s) - 8 GiB

Cancel

Launch instance

[Review commands](#)

The summary of our command.

Using the AWS console connect the instance to the load balancer. Watch the status of the instance go from Out of Service to In Service.

Register targets

Select instances, specify ports, and add the instances to the list of pending targets. Repeat to add additional combinations of instances and ports to the list of pending targets. Once you are satisfied with your selections, click Register pending targets.

Available instances (2)

Filter instances

< 1 >

| <input type="checkbox"/> | Instance ID | Name | State | Security groups | Zone | Private IPv4 address | Subnet ID |
|--------------------------|---------------------|--------------------|---------|----------------------|------------|----------------------|--------------------------|
| <input type="checkbox"/> | i-01f5ba21b75f4acbc | GrR_Auberson_Image | Running | GrR_Auberson_default | us-east-1d | 172.31.19.193 | subnet-0a2ab628966261f50 |
| <input type="checkbox"/> | i-0e422f6114fba9dd | GrR_Auberson | Running | GrR_Auberson_default | us-east-1d | 172.31.17.184 | subnet-0a2ab628966261f50 |

0 selected

Ports for the selected instances

Ports for routing traffic to the selected instances.

80

1-65535 (separate multiple ports with commas)

Include as pending below

2 selections are now pending below. Include more or register targets when ready.

Review targets

Targets (2)

Filter targets

Show only pending

Remove all pending

< 1 >

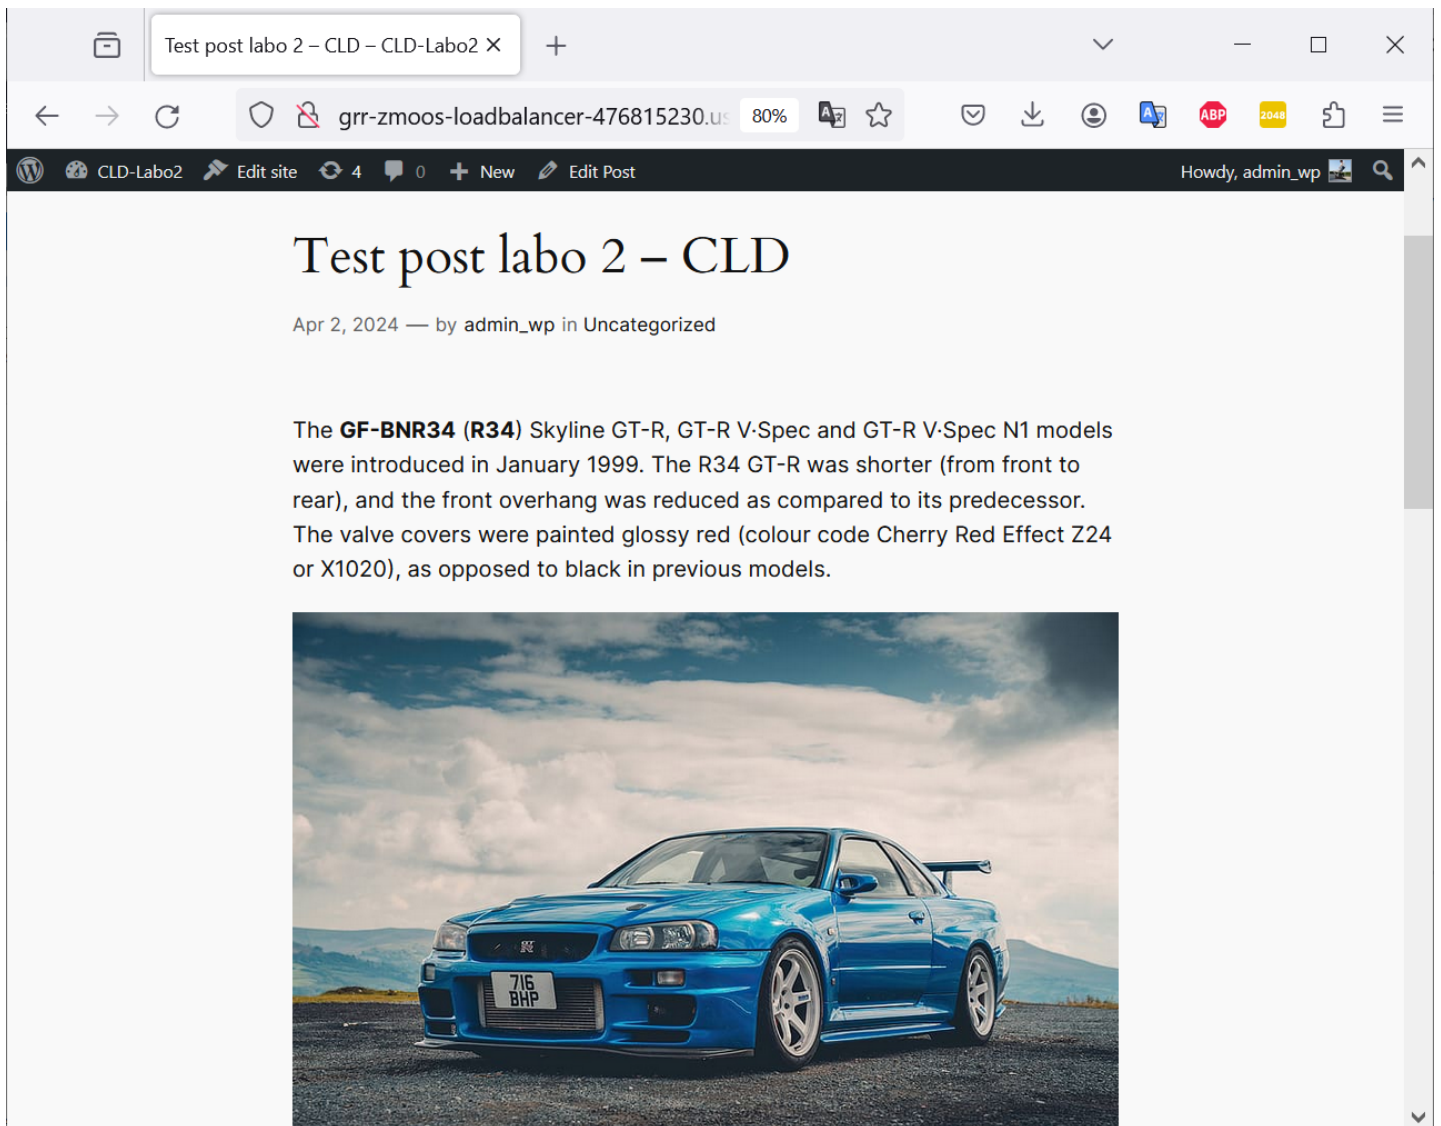
| Instance ID | Name | Port | State | Security groups | Zone | Private IPv4 address | Subnet ID | Launch time |
|---------------------|--------------------|------|---------|----------------------|------------|----------------------|--------------------------|----------------------------------|
| i-01f5ba21b75f4acbc | GrR_Auberson_Image | 80 | Running | GrR_Auberson_default | us-east-1d | 172.31.19.193 | subnet-0a2ab628966261f50 | April 8, 2024, 18:39 (UTC+02:00) |
| i-0e422f6114fba9dd | GrR_Auberson | 80 | Running | GrR_Auberson_default | us-east-1d | 172.31.17.184 | subnet-0a2ab628966261f50 | April 8, 2024, 15:29 (UTC+02:00) |

Adding instance in load balancer

Using any of the instances, run the wp search and replace tool to replace the old IP address with the load balancer’s DNS name in the database.

```
php wp-cli.phar search-replace '52.54.125.85' 'GrR-Auberson-LoadBalancer-13'
```

Make sure that you can access the Wordpress post using the load balancer’s DNS name.



Access to my page create with loadbalancer DNS name

Draw a diagram of the setup you have created showing the components (instances, database, load balancer, client) and how they are connected. Include the security groups as well. Make sure to show every time a packet is filtered.

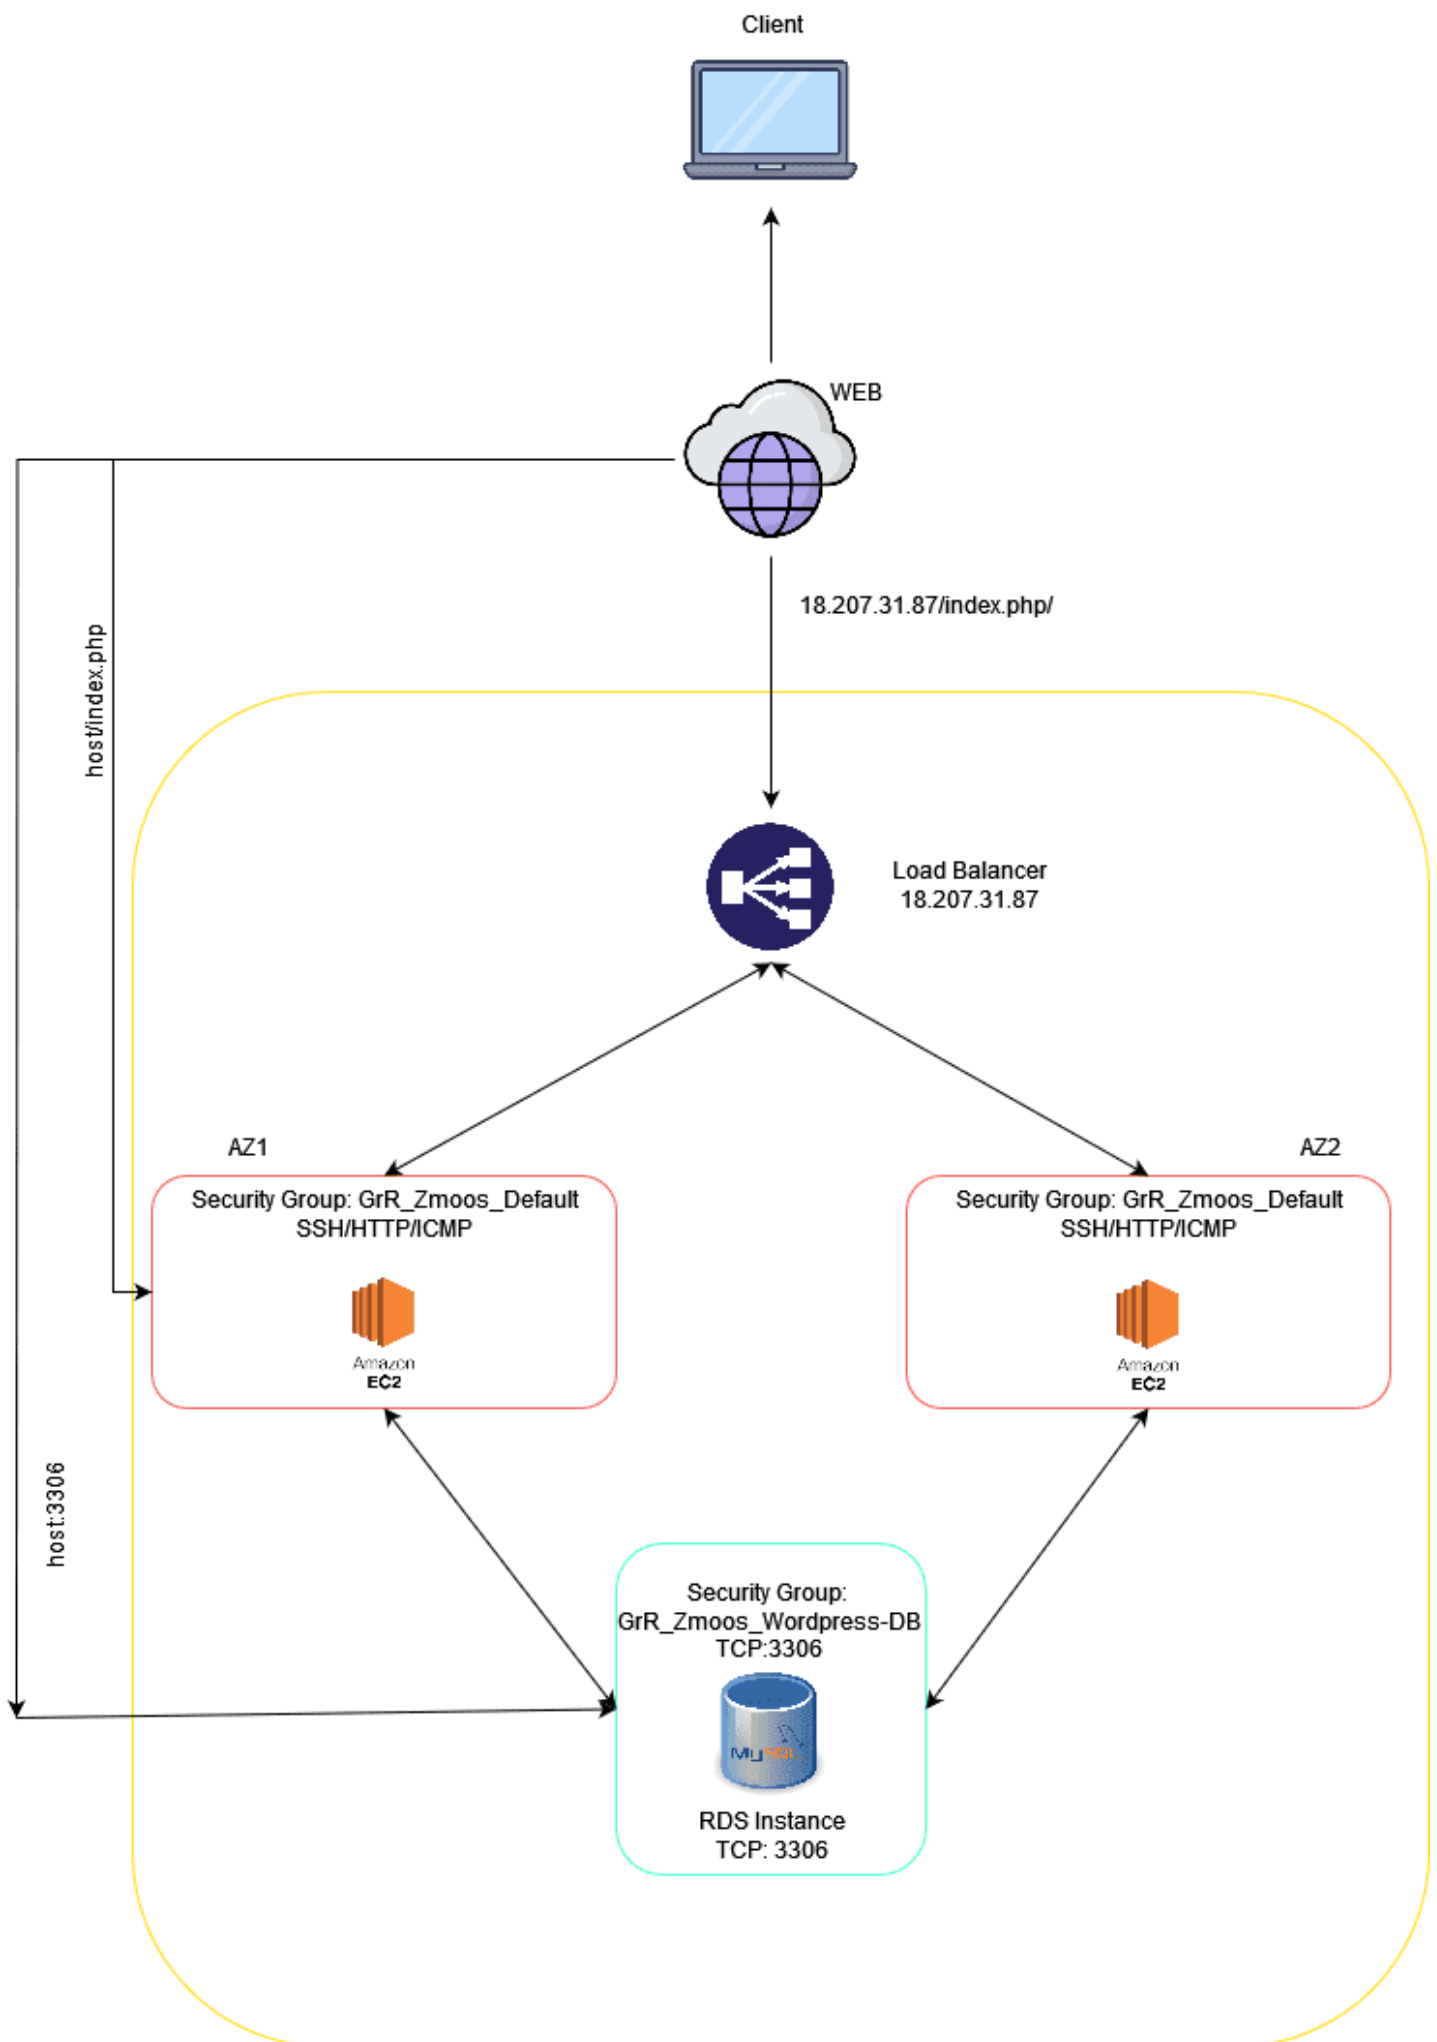


Diagram of the setup

Diagram of the setup

Calculate the monthly cost of this setup. You can ignore traffic costs.

Estimate summary [Info](#)

Upfront cost
0.00 USD

Monthly cost
62.95 USD

Total 12 months cost
755.40 USD
Includes upfront cost

Getting Started with AWS

[Get started for free](#)[Contact Sales](#)

My Estimate

[Duplicate](#)[Delete](#)[Move to](#)[Create group](#)[Add support](#)[Add service](#)

< 1 >

| <input type="checkbox"/> | Service Name | Status | Upfron... | Monthly cost | Descrip... | Region | Config Summary |
|--------------------------|------------------------|--------|-----------|--------------|------------|---------------------|-------------------------|
| <input type="checkbox"/> | Amazon RDS for MySQL | - | 0.00 USD | 29.42 USD | - | US East (N. Virg... | Storage for each RD... |
| <input type="checkbox"/> | Amazon EC2 | - | 0.00 USD | 16.94 USD | - | US East (N. Virg... | Tenancy (Shared Inst... |
| <input type="checkbox"/> | Elastic Load Balancing | - | 0.00 USD | 16.59 USD | - | US East (N. Virg... | Number of Applicati... |

- RDS with a 100%utilized/Month
 - EC2 constant usage with 100%utilized/Month
 - Load balancing with 20 GB/month
- The monthly cost is 62.95 USD.

TASK 5B: DELETE AND RE-CREATE THE LOAD BALANCER USING THE COMMAND LINE INTERFACE

Put the commands to delete the load balancer, re-create the load balancer and re-create the listener into the report.

Commands:

```
//For deleting the loadbalancer
aws elbv2 delete-load-balancer --load-balancer-arn arn:aws:elasticloadbalancing:us-east-1:123456789012:loadbalancer/app/my-load-balancer/50dc6c49-6ba0-40a6-83d0-d6a21b061b6c

//For create a loadbalancer
aws elbv2 create-load-balancer --name GrR-Zmoos-LoadBalancer --subnets subnet-12345678


//For create listener
aws elbv2 create-listener --load-balancer-arn arn:aws:elasticloadbalancing:us-east-1:123456789012:loadbalancer/app/my-load-balancer/50dc6c49-6ba0-40a6-83d0-d6a21b061b6c --protocol HTTPS --port 443 --default-target-groups tg-12345678
```

TASK 6: TEST THE DISTRIBUTED APPLICATION

Document your observations. Include reports and graphs of the load testing tool and the AWS console monitoring output.

After running Vegeta tool once:

Master instance:



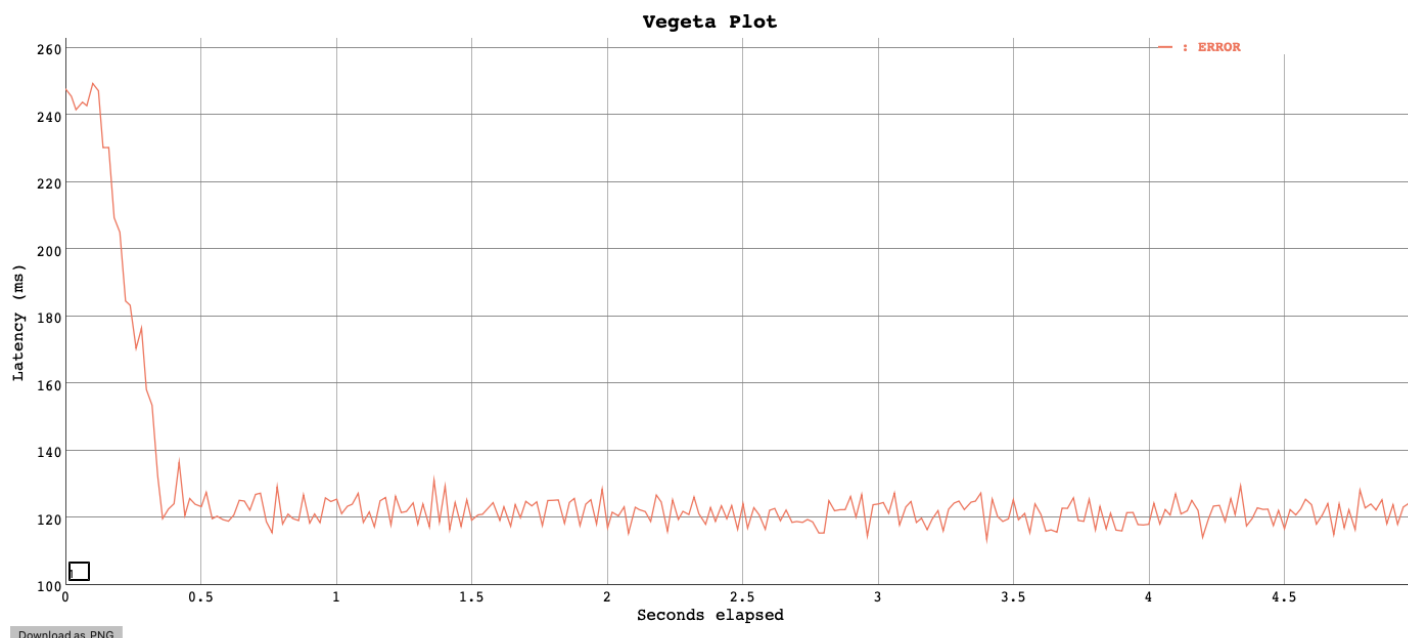
```
ubuntu@ip-172-31-26-63: ~  
/1.1" 200 12168 "-" "Go-http-client/1.1"  
172.31.11.220 - - [08/Apr/2024:14:50:52 +0000] "GET / HTTP/1.1" 200 3423 "-" "ELB-HealthChecker/2.0"  
172.31.11.220 - - [08/Apr/2024:14:50:50 +0000] "GET /index.php/sample-page/ HTTP/1.1" 200 12168 "-" "Go-http-client/1.1"  
172.31.11.220 - - [08/Apr/2024:14:50:50 +0000] "GET /index.php/sample-page/ HTTP/1.1" 200 12168 "-" "Go-http-client/1.1"  
172.31.11.220 - - [08/Apr/2024:14:53:20 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 501 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:53:20 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 501 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:53:20 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 0 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:53:20 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 501 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:53:20 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 0 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:54:10 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 501 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:54:11 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 501 "-" "-"  
172.31.11.220 - - [08/Apr/2024:14:54:04 +0000] "GET /index.php/sample-page/ HTTP/1.1" 408 501 "-" "-"
```

Second instance:




Vegeta plot:

Duration: 60s



Duration: 120s



Observations:

Before we started testing with Vegeta, when we ran the command `sudo tail -F /var/log/apache2/access.log` on the second instance, the log file was empty because for the moment all incoming requests were directed to the master instance.

However, this behaviour changed during tests with Vegeta! Obviously, this was due to the load balancer redirecting some of the requests to the second instance in order to avoid overloading the first.

We can clearly see from the two AWS graphs that the load on the first instance decreased as soon as the second instance started processing requests too.

This indicates that our load balancer is working properly.

When you resolve the DNS name of the load balancer into IP addresses what do you see? Explain.

```
C:\Users\Dell>nslookup GrR-Zmoos-LoadBalancer-476815230.us-east-1.elb.amazonaws.com
Server: fritz.box
Address: fd00::de39:6fff:fe41:c1ad

Non-authoritative answer:
Name: GrR-Zmoos-LoadBalancer-476815230.us-east-1.elb.amazonaws.com
Addresses: 3.223.194.21
           52.202.25.12

C:\Users\Dell>
```

We can see 2 different IP addresses associated to the DNS (A record).

Did this test really test the load balancing mechanism? What are the limitations of this simple test? What would be necessary to do realistic testing?

Our current setup effectively evaluates the load balancing mechanism of our Application Load Balancer by distributing traffic among EC2 instances, it has its limitations. This basic test overlooks the diversity of traffic patterns in real-world scenarios. To (http://scenarios.To) fully test the system's performance, we must conduct a more varied set of tests that accurately simulate real-world scenarios

To perform realistic test, we could use different types of requests, varying the frequency and simulating different user. Furthermore, we could consider testing the system under different conditions, such as the peak hours traffic.